

Data Cleaning Code

November 15, 2021

Kiet Nguyen

ID: 001601720

Email: kngu179@wgu.edu

Part I: Research Question

A. Question to address

What variables in the dataset can help predict which customers are most likely to discontinue service?

B. Variables description

```
import pandas as pd

# Import CSV file to Pythondf.
churn = pd.read_csv('churn_raw_data.csv')

churn.shape          # Return the rows x columns of dataset
churn.info()         # Print summary of dataframe
```

Running the above code, the output shows that the dataset contains 52 variables and 10,000 records. The variable we are interested in is **Churn**, whether the customer discontinued service within the last month. The dataset also includes services that each customer signed up for, customer account information, and customer demographics. The data types of the variables are **int64**, **object**, and **float64**.

Part II: Data Cleaning Plan

C1. Plan to identify anomalies

My plan to find the anomalies include the following steps:

- Learn about the dataset and its characteristics.
- Explore measures of central tendency (mean, median, mode) and potential outliers.
- Summarize and visualize missing values for possible patterns.
- Impute missing data.

C2. Approach to assess data quality

To assess the data quality, I need to answer the following questions:

- Is the data accurate?
- How much data is missing?

- Is the data in the correct format?
- Are there duplicates?

The most difficult question to answer is about data accuracy since I do not have a baseline to compare this dataset against. The other questions can be answered through exploration of the data. Finding out how much data is missing will help me to avoid poor estimations and analysis based on incomplete data (Donthi, 2021). Incorrect formatting is another issue that can cause problems with processing data. Having a lot of duplicates can also skew the results of the analysis.

C3. Language and libraries

For the analysis of the data, I will use Python. It is a powerful general purpose programming language with an easy to read syntax. It also provides specialized libraries for working with data science projects.

A. Nehme (2021) recommended the following Python libraries for data cleaning:

- `pandas` - data manipulation and analysis.
- `numpy` - math functions and multidimensional arrays and matrices support.
- `matplotlib` - data visualization.
- `missingno` - missing data visualization.
- `scikit-learn` - machine learning models.

C4. Code to identify anomalies

```
[26]: import pandas as pd
import missingno as msno
import matplotlib.pyplot as plt

# Import CSV to dataframe
churn = pd.read_csv('churn_raw_data.csv')
```

```
[27]: # Return the rows x columns of dataset
churn.shape
```

```
[27]: (10000, 52)
```

```
[28]: # Print summary of dataframe
churn.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 52 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0            10000 non-null  int64
1   CaseOrder             10000 non-null  int64
2   Customer_id           10000 non-null  object
3   Interaction            10000 non-null  object
4   City                  10000 non-null  object
5   State                 10000 non-null  object
```

6	County	10000	non-null	object
7	Zip	10000	non-null	int64
8	Lat	10000	non-null	float64
9	Lng	10000	non-null	float64
10	Population	10000	non-null	int64
11	Area	10000	non-null	object
12	Timezone	10000	non-null	object
13	Job	10000	non-null	object
14	Children	7505	non-null	float64
15	Age	7525	non-null	float64
16	Education	10000	non-null	object
17	Employment	10000	non-null	object
18	Income	7510	non-null	float64
19	Marital	10000	non-null	object
20	Gender	10000	non-null	object
21	Churn	10000	non-null	object
22	Outage_sec_perweek	10000	non-null	float64
23	Email	10000	non-null	int64
24	Contacts	10000	non-null	int64
25	Yearly equip_failure	10000	non-null	int64
26	Techie	7523	non-null	object
27	Contract	10000	non-null	object
28	Port_modem	10000	non-null	object
29	Tablet	10000	non-null	object
30	InternetService	10000	non-null	object
31	Phone	8974	non-null	object
32	Multiple	10000	non-null	object
33	OnlineSecurity	10000	non-null	object
34	OnlineBackup	10000	non-null	object
35	DeviceProtection	10000	non-null	object
36	TechSupport	9009	non-null	object
37	StreamingTV	10000	non-null	object
38	StreamingMovies	10000	non-null	object
39	PaperlessBilling	10000	non-null	object
40	PaymentMethod	10000	non-null	object
41	Tenure	9069	non-null	float64
42	MonthlyCharge	10000	non-null	float64
43	Bandwidth_GB_Year	8979	non-null	float64
44	item1	10000	non-null	int64
45	item2	10000	non-null	int64
46	item3	10000	non-null	int64
47	item4	10000	non-null	int64
48	item5	10000	non-null	int64
49	item6	10000	non-null	int64
50	item7	10000	non-null	int64
51	item8	10000	non-null	int64

dtypes: float64(9), int64(15), object(28)

memory usage: 4.0+ MB

```
[29]: # See first 5 rows
churn.head()
```

```
[29]: Unnamed: 0 CaseOrder Customer_id Interaction \
0      1      1      K409198 aa90260b-4141-4a24-8e36-b04ce1f4f77b
1      2      2      S120509 fb76459f-c047-4a9d-8af9-e0f7d4ac2524
2      3      3      K191035 344d114c-3736-4be5-98f7-c72c281e2d35
3      4      4      D90850 abfa2b40-2d43-4994-b15a-989b8c79e311
4      5      5      K662701 68a861fd-0d20-4e51-a587-8a90407ee574

      City State      County Zip Lat Lng ... \
0 Point Baker AK Prince of Wales-Hyder 99927 56.25100 -133.37571 ...
1 West Branch MI Ogemaw 48661 44.32893 -84.24080 ...
2 Yamhill OR Yamhill 97148 45.35589 -123.24657 ...
3 Del Mar CA San Diego 92014 32.96687 -117.24798 ...
4 Needville TX Fort Bend 77461 29.38012 -95.80673 ...

      MonthlyCharge Bandwidth_GB_Year item1 item2 item3 item4 item5 item6 \
0 171.449762 904.536110 5 5 5 3 4 4
1 242.948015 800.982766 3 4 3 3 4 3
2 159.440398 2054.706961 4 4 2 4 4 3
3 120.249493 2164.579412 4 4 4 2 5 4
4 150.761216 271.493436 4 4 4 3 4 4

      item7 item8
0 3 4
1 4 4
2 3 3
3 3 3
4 4 5

[5 rows x 52 columns]
```

```
[30]: # Get statistics of numeric columns
churn.describe()
```

```
[30]: Unnamed: 0 CaseOrder Zip Lat Lng \
count 10000.00000 10000.00000 10000.000000 10000.000000 10000.000000
mean 5000.50000 5000.50000 49153.319600 38.757567 -90.782536
std 2886.89568 2886.89568 27532.196108 5.437389 15.156142
min 1.00000 1.00000 601.000000 17.966120 -171.688150
25% 2500.75000 2500.75000 26292.500000 35.341828 -97.082812
50% 5000.50000 5000.50000 48869.500000 39.395800 -87.918800
75% 7500.25000 7500.25000 71866.500000 42.106908 -80.088745
max 10000.00000 10000.00000 99929.000000 70.640660 -65.667850

      Population Children Age Income \
```

count	10000.000000	7505.000000	7525.000000	7510.000000
mean	9756.562400	2.095936	53.275748	39936.762226
std	14432.698671	2.154758	20.753928	28358.469482
min	0.000000	0.000000	18.000000	740.660000
25%	738.000000	0.000000	35.000000	19285.522500
50%	2910.500000	1.000000	53.000000	33186.785000
75%	13168.000000	3.000000	71.000000	53472.395000
max	111850.000000	10.000000	89.000000	258900.700000

	Outage_sec_perweek	...	MonthlyCharge	Bandwidth_GB_Year	\
count	10000.000000	...	10000.000000	8979.000000	
mean	11.452955	...	174.076305	3398.842752	
std	7.025921	...	43.335473	2187.396807	
min	-1.348571	...	77.505230	155.506715	
25%	8.054362	...	141.071078	1234.110529	
50%	10.202896	...	169.915400	3382.424000	
75%	12.487644	...	203.777441	5587.096500	
max	47.049280	...	315.878600	7158.982000	

	item1	item2	item3	item4	item5	\
count	10000.000000	10000.000000	10000.000000	10000.000000	10000.000000	
mean	3.490800	3.505100	3.487000	3.497500	3.492900	
std	1.037797	1.034641	1.027977	1.025816	1.024819	
min	1.000000	1.000000	1.000000	1.000000	1.000000	
25%	3.000000	3.000000	3.000000	3.000000	3.000000	
50%	3.000000	4.000000	3.000000	3.000000	3.000000	
75%	4.000000	4.000000	4.000000	4.000000	4.000000	
max	7.000000	7.000000	8.000000	7.000000	7.000000	

	item6	item7	item8
count	10000.000000	10000.000000	10000.000000
mean	3.497300	3.509500	3.495600
std	1.033586	1.028502	1.028633
min	1.000000	1.000000	1.000000
25%	3.000000	3.000000	3.000000
50%	3.000000	4.000000	3.000000
75%	4.000000	4.000000	4.000000
max	8.000000	7.000000	8.000000

[8 rows x 24 columns]

```
[31]: # Get statistics of non-numeric columns
churn.describe(include=object)
```

```
[31]:      Customer_id      Interaction      City      State \
count      10000      10000      10000      10000
unique      10000      10000      6058      52
```

top	K409198	aa90260b-4141-4a24-8e36-b04ce1f4f77b	Houston	TX
freq	1		1	34 603

	County	Area	Timezone	Job	\
count	10000	10000	10000	10000	
unique	1620	3	25	639	
top	Washington	Suburban	America/New_York	Occupational psychologist	
freq	111	3346	4072	30	

	Education	Employment	... Phone	Multiple	\
count	10000	10000	... 8974	10000	
unique	12	5	... 2	2	
top	Regular High School Diploma	Full Time	... Yes	No	
freq	2421	5992	... 8128	5392	

	OnlineSecurity	OnlineBackup	DeviceProtection	TechSupport	StreamingTV	\
count	10000	10000	10000	9009	10000	
unique	2	2	2	2	2	
top	No	No	No	No	No	
freq	6424	5494	5614	5635	5071	

	StreamingMovies	PaperlessBilling	PaymentMethod
count	10000	10000	10000
unique	2	2	4
top	No	Yes	Electronic Check
freq	5110	5882	3398

[4 rows x 28 columns]

```
[32]: # Check if there is any duplicate across all columns
duplicates = churn.duplicated().any()

# No duplicate found
duplicates
```

[32]: False

```
[33]: # Print unique values of each categorical column
# Print total if there are more than 15 unique values
for column in churn:
    if churn[column].dtype == object:
        print(column)
        uniques = churn[column].unique()
        if len(uniques) > 15:
            print(f'- {len(uniques)}')
        else:
            print(f'- {churn[column].unique()}')
```

Customer_id
 - 10000
 Interaction
 - 10000
 City
 - 6058
 State
 - 52
 County
 - 1620
 Area
 - ['Urban' 'Suburban' 'Rural']
 Timezone
 - 25
 Job
 - 639
 Education
 - ["Master's Degree" 'Regular High School Diploma' 'Doctorate Degree'
 'No Schooling Completed' "Associate's Degree" "Bachelor's Degree"
 'Some College, Less than 1 Year' 'GED or Alternative Credential'
 'Some College, 1 or More Years, No Degree'
 '9th Grade to 12th Grade, No Diploma' 'Nursery School to 8th Grade'
 'Professional School Degree']
 Employment
 - ['Part Time' 'Retired' 'Student' 'Full Time' 'Unemployed']
 Marital
 - ['Widowed' 'Married' 'Separated' 'Never Married' 'Divorced']
 Gender
 - ['Male' 'Female' 'Prefer not to answer']
 Churn
 - ['No' 'Yes']
 Techie
 - ['No' 'Yes' nan]
 Contract
 - ['One year' 'Month-to-month' 'Two Year']
 Port_modem
 - ['Yes' 'No']
 Tablet
 - ['Yes' 'No']
 InternetService
 - ['Fiber Optic' 'DSL' 'None']
 Phone
 - ['Yes' 'No' nan]
 Multiple
 - ['No' 'Yes']
 OnlineSecurity
 - ['Yes' 'No']
 OnlineBackup

```

- ['Yes' 'No']
DeviceProtection
- ['No' 'Yes']
TechSupport
- ['No' 'Yes' nan]
StreamingTV
- ['No' 'Yes']
StreamingMovies
- ['Yes' 'No']
PaperlessBilling
- ['Yes' 'No']
PaymentMethod
- ['Credit Card (automatic)' 'Bank Transfer(automatic)' 'Mailed Check'
  'Electronic Check']

```

```

[34]: # Explore columns with missing values similar to the 'nan' results above
      churn_null = churn.isnull()
      null_sums = churn_null.sum()

      # Add missing column names to a list print out missing percentage
      print('Total and percentage of missing values per column')
      cols_missing = []
      for column, missing in null_sums.items():
          if missing > 0:
              print(f'{column}: {missing} - {round((missing / 10000) * 100, )}%')
              cols_missing.append(column)

```

```

Total and percentage of missing values per column
Children: 2495 - 25%
Age: 2475 - 25%
Income: 2490 - 25%
Techie: 2477 - 25%
Phone: 1026 - 10%
TechSupport: 991 - 10%
Tenure: 931 - 9%
Bandwidth_GB_Year: 1021 - 10%


```

```

[35]: # Copy missing columns to new dataframe
      churn_missing = churn[cols_missing].copy()

      # Visualize missing data
      msno.bar(churn_missing, figsize=(12, 12))
      plt.show()

```

```
[36]: """
      Author: Brad Solomon
      Date: 2021
      Title: Python Histogram Plotting: NumPy, Matplotlib, Pandas & Seaborn
      Type: Source code
      Web: https://realpython.com/python-histograms/
      """

      # Plot the histograms of interesting numeric columns to see patterns
```

```
churn[['Children', 'Age', 'Income', 'Outage_sec_perweek', 'Contacts',
      ↪ 'Yearly_equip_failure', 'Tenure', 'MonthlyCharge', 'Bandwidth_GB_Year']].
      ↪ hist(figsize=(12, 12),
```

```
      ↪
```

```
      ↪ ec='black',
```

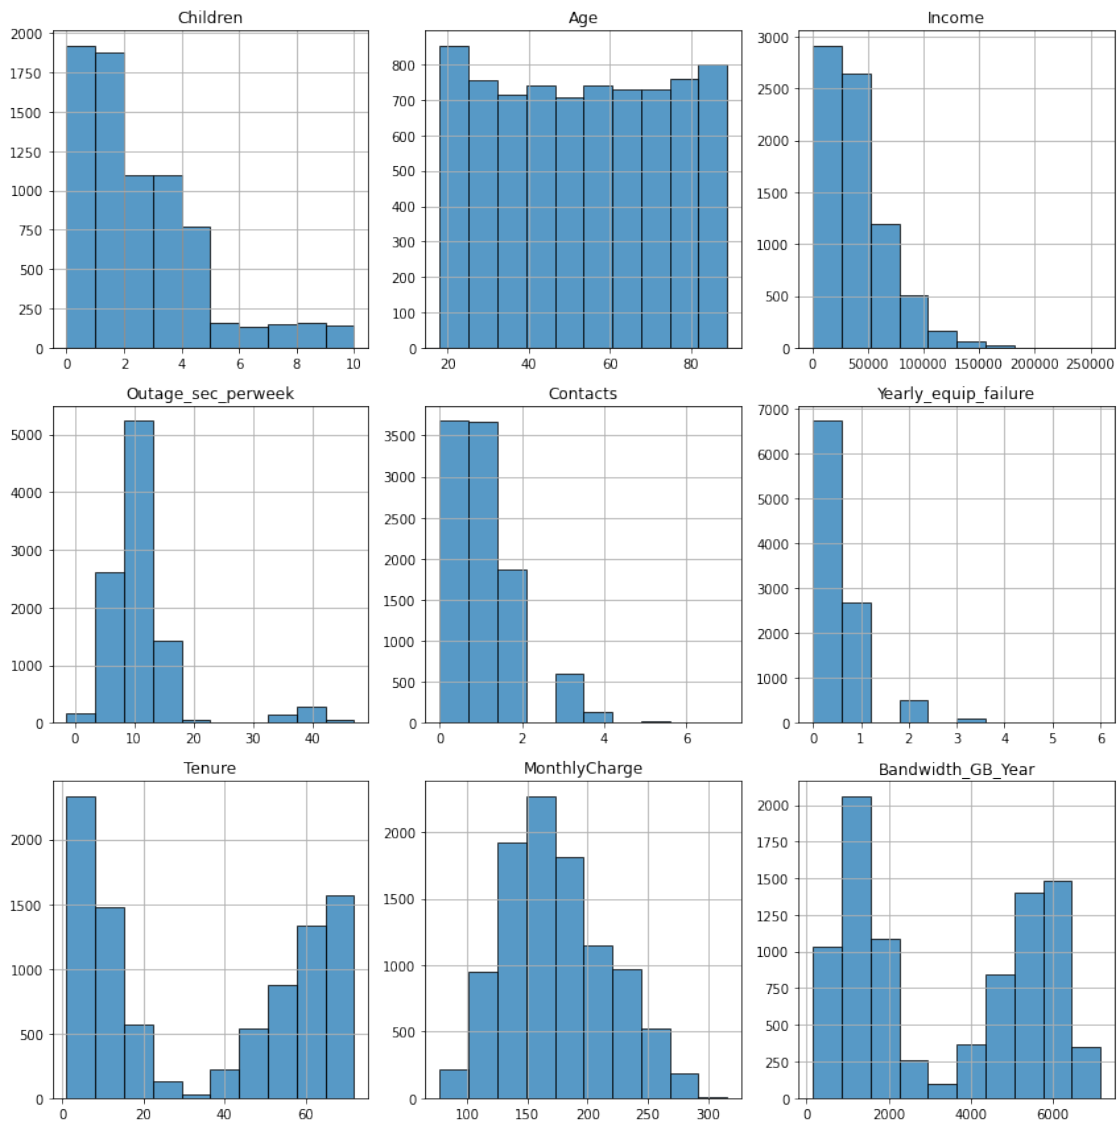
```
      ↪
```

```
      ↪ 75)
```

```
plt.tight_layout()
```


```
plt.show()
```

alpha=0.



```
[37]: # Plot survey responses
churn[[f'item{i}' for i in range(1, 9)]].hist(figsize=(12, 12),
                                              ec='black',
                                              alpha=0.75,
                                              color='green',
                                              bins=[1, 2, 3, 4, 5, 6, 7, 8])

plt.tight_layout()
plt.show()
```




```
[38]: """
Author: Python Pool
Date: 2021
Title: Matplotlib Boxplot With Customization in Python
"""
```

Type: Source Code

Web: <https://www.pythonpool.com/matplotlib-boxplot/>

```
"""  
  
# Verify columns with outliers using boxplots  
fig, ax = plt.subplots(3, 1, figsize=(15, 10))  
for index, column in enumerate(['Outage_sec_perweek', 'Contacts',  
    ↳ 'Yearly equip_failure']):  
    ax[index].boxplot(  
        churn[column],  
        vert=False,  
        patch_artist=True,  
        flierprops=dict(markerfacecolor='red', markersize=10, linestyle='none',  
    ↳ alpha=0.6)  
    )  
    ax[index].set_title(column)  
plt.show()
```



Part III: Data Cleaning

D1. Describe findings

There are several anomalies in the dataset:

- Columns `Unnamed: 0` and `CaseOrder` are the same.
- Unclear column names for survey responses.
- Negative values in `Outage_sec_perweek`, where the minimum should only be zero.
- Multiple columns are missing from 9% to 25% of their data.
- Based on the histograms and boxplots, there are outliers for `Outage_sec_perweek`, `Contacts`, and `Yearly_equip_failure`.

D2. Mitigation methods

My approach for dealing with the anomalies will be:

- Drop `Unnamed: 0` column since it's a duplicate of `CaseOrder`.
- Rename survey responses to something easier to understand.
- Change the negative values in `Outage_sec_perweek` to 0 to enforce proper range (Nehme, 2021).
- Use K-Nearest Neighbor algorithm to predict missing values based on values of neighboring columns (Donthi, 2021).
- Leave the outliers alone since they might provide important insights for later analysis.

D3. Outcome summary

The outcome of each data cleaning step is:

- `Unnamed: 0` column dropped from the dataframe.
- Survey items are easier to understand.
- The minimum for `Outage_sec_perweek` is set to 0.
- Missing values are imputed using KNN algorithm. Categorical values are ordinally encoded before imputation.

D4. Code for mitigation

```
[39]: # Drop 'Unnamed: 0' column
churn = churn.drop('Unnamed: 0', axis=1)
```

```
[40]: # Rename survey responses columns
churn.rename(columns={
    'item1': 'SurveyResponse',
    'item2': 'SurveyFixes',
    'item3': 'SurveyReplacements',
    'item4': 'SurveyReliability',
    'item5': 'SurveyOptions',
    'item6': 'SurveyRespect',
    'item7': 'SurveyCourteous',
    'item8': 'SurveyListening'
}, inplace=True)
```

```
[41]: # Set negative values in Outage_sec_perweek to 0
churn.loc[churn['Outage_sec_perweek'] < 0, 'Outage_sec_perweek'] = 0
```

```
[42]: import numpy as np
from sklearn.impute import KNNImputer
```

```

# Split missing columns into numeric and categorical
num_missing = churn_missing[['Children', 'Age', 'Income', 'Tenure',
    ↳ 'Bandwidth_GB_Year']].copy()
cat_missing = churn_missing[['Techie', 'Phone', 'TechSupport']].copy()

# Impute numeric columns
knn = KNNImputer()
num_missing.iloc[:, :] = knn.fit_transform(num_missing)

num_missing.describe()

```

```

[42]:

```

	Children	Age	Income	Tenure \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	2.045900	52.99086	39976.469551	34.539410
std	1.917549	18.62637	25328.964531	25.539056
min	0.000000	18.00000	740.660000	1.000259
25%	1.000000	39.20000	22838.322500	8.574026
50%	1.800000	52.40000	35629.102000	34.369232
75%	3.000000	67.00000	50947.115000	60.367002
max	10.000000	89.00000	258900.700000	71.999280

	Bandwidth_GB_Year
count	10000.000000
mean	3388.508002
std	2108.192078
min	155.506715
25%	1290.549139
50%	3349.915681
75%	5476.439000
max	7158.982000

```

[43]: """
Author: Suraj Donthi
Date: 2021
Title: Imputing categorical values
Type: Source code
Web: https://campus.datacamp.com/courses/dealing-with-missing-data-in-python/
    ↳ advanced-imputation-techniques?ex=4
"""

from sklearn.preprocessing import OrdinalEncoder

# Perform ordinal encoding for missing categorical columns
ordinal_dict = {}
for column in cat_missing:
    # Create ordinal encoder for column

```

```

ordinal_dict[column] = OrdinalEncoder()
col = cat_missing[column]

# Select non-null values and encode them
col_not_null = col[col.notnull()]
reshaped_vals = col_not_null.values.reshape(-1, 1)
encoded_vals = ordinal_dict[column].fit_transform(reshaped_vals)

# Store the encoded values back in the column
cat_missing.loc[col.notnull(), column] = np.squeeze(encoded_vals)

# Impute missing values
cat_missing.iloc[:, :] = np.round(knn.fit_transform(cat_missing))

# Convert encoded columns back to categorical
for column in cat_missing:
    reshaped = cat_missing[column].values.reshape(-1, 1)
    cat_missing[column] = ordinal_dict[column].inverse_transform(reshaped)

cat_missing.describe()

```

```

[43]:
      Techie  Phone TechSupport
count  10000  10000      10000
unique      2      2          2
top        No    Yes          No
freq      8743  9154      6413

```

D5. Copy of clean data

```

[44]: # Create copy of churn data
churn_clean = churn.copy(deep=True)

# Replace missing columns with imputed ones
for column in churn_missing:
    if column in num_missing:
        churn_clean[column] = num_missing[column]
    elif column in cat_missing:
        churn_clean[column] = cat_missing[column]

churn_clean.to_csv('churn_clean.csv')

```

```

[45]: # Check the cleaned data by importing and display summary of imputed columns
cleaned_data = pd.read_csv('churn_clean.csv')

cleaned_data[['Children', 'Age', 'Income', 'Tenure', 'Bandwidth_GB_Year',
               'Techie', 'Phone', 'TechSupport']].info()

```

```
<class 'pandas.core.frame.DataFrame'>
```

```

RangeIndex: 10000 entries, 0 to 9999
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   Children              10000 non-null  float64
 1   Age                   10000 non-null  float64
 2   Income                10000 non-null  float64
 3   Tenure                10000 non-null  float64
 4   Bandwidth_GB_Year    10000 non-null  float64
 5   Techie                10000 non-null  object  
 6   Phone                 10000 non-null  object  
 7   TechSupport           10000 non-null  object  
dtypes: float64(5), object(3)
memory usage: 625.1+ KB

```

D6. Limitations of data cleaning

Without knowing where the data came from or how it was collected, I have to assume some uncertainty when working with it. While KNN algorithm is better than using the mean or median for imputing values (Donthi, 2021), , it does not replace real data. By imputing the values, I might lose some insight when performing analysis with the imputed values. To properly clean the data, I would need help from a domain expert, someone who understands churn in both the company and the industry. Since I'm not familiar with the industry, my data cleaning process might be very different than someone who understands the industry well.

D7. Effects of limitation on research question

Given that the dataset was missing a lot of data, the conclusion drawn from this data needs to be scrutinized before decisions are made from imputed values (Nehme, 2021). Predicting which customers are at high risk of churn require some degree of caution. It can be very costly for the company to predict the wrong segment of customers to retain, leading to ineffective marketing. It's also possible for the company to focus on the wrong metrics for customer retention based on erroneous data. There are also some variables in the data that need more consideration since their correlations with churn are unclear at the moment.

```

[46]: # Create dataframe with only numeric columns for PCA
churn_cp = churn_clean.select_dtypes(include=['int64', 'float64'])

# Drop columns that are not meaningful
churn_cp = churn_cp.drop(['CaseOrder', 'Zip', 'Lat', 'Lng'], axis=1)

churn_cp.head()

```

```

[46]:   Population  Children  Age    Income  Outage_sec_perweek  Email  Contacts  \
0          38         1.2  68.0  28561.990           6.972566     10         0
1       10446         1.0  27.0  21704.770          12.014541     12         0
2        3735         4.0  50.0  45759.278          10.245616      9         0
3       13863         1.0  48.0  18925.230          15.206193     15         2
4       11352         0.0  83.0   40074.190           8.960316     16         2

```


	Yearly_equip_failure	Tenure	MonthlyCharge	Bandwidth_GB_Year	\
0	1	6.795513	171.449762	904.536110	
1	1	1.156681	242.948015	800.982766	
2	1	15.754144	159.440398	2054.706961	
3	0	17.087227	120.249493	2164.579412	
4	1	1.670972	150.761216	271.493436	

	SurveyResponse	SurveyFixes	SurveyReplacements	SurveyReliability	\
0	5	5	5	3	
1	3	4	3	3	
2	4	4	2	4	
3	4	4	4	2	
4	4	4	4	3	

	SurveyOptions	SurveyRespect	SurveyCourteous	SurveyListening
0	4	4	3	4
1	4	3	4	4
2	4	3	3	3
3	5	4	3	3
4	4	4	4	5

```
[47]: """
Author: Keiona Middleton
Date: 2021
Title: How to perform PCA in Python
Type: Source code
Web: https://cgp-oex.wgu.edu/lti\_provider/courses/course-v1:WGUx+D206x+2021\_T3/
↳ block-v1:
↳ WGUx+D206x+2021_T3+type@vertical+block@3f0422d47d8b4a3eaec25de3bced8bf8
"""

from sklearn.decomposition import PCA

# Normalize data
churn_normalized = (churn_cp - churn_cp.mean()) / churn_cp.std()


# Extract all components for analysis
pca = PCA(n_components=churn_cp.shape[1])
pca.fit(churn_normalized)

# Converts 19 variables into dataset of 19 components
pca_cols = [f'PC{i}' for i in range(1, 20)] # Create list of component_
↳ names 'PC1' to 'PC19'
churn_pca = pd.DataFrame(pca.transform(churn_normalized),
                        columns=pca_cols)
```

```

# Scree plot of PCA with each additional component variance
pc_values = np.arange(pca.n_components_) + 1
plt.figure(figsize=(12, 8))
plt.plot(pc_values, pca.explained_variance_ratio_, 'o-', color='blue')
plt.xlabel('Number of components')
plt.ylabel('Explained variance')
plt.show()

```




```

[48]: # Extract eigenvalues from dataset
cov_matrix = np.dot(churn_normalized.T, churn_normalized) / churn_cp.shape[0]
eigenvalues = [np.dot(eigenvector.T, np.dot(cov_matrix, eigenvector)) for
               eigenvector in pca.components_]

# Plot eigenvalues
plt.figure(figsize=(12, 8))
plt.plot(pc_values, eigenvalues, 'o-', color='green')
plt.xlabel('Number of components')
plt.ylabel('Eigenvalue')
plt.show()

```



```
[49]: # Table output of loadings for the components
loadings = pd.DataFrame(pca.components_.T,
                        columns=pca_cols,
                        index=churn_cp.columns)

loadings
```

```
[49]:
```

	PC1	PC2	PC3	PC4	PC5	\
Population	-0.002126	-0.000453	0.015690	-0.044868	-0.460027	
Children	0.003095	0.012553	0.004845	0.011792	0.421564	
Age	0.005289	-0.008653	-0.014067	-0.056497	0.051788	
Income	0.000406	0.026147	0.014560	-0.006958	0.114120	
Outage_sec_perweek	-0.013174	0.016744	-0.047773	0.703619	0.010886	
Email	0.008630	-0.022362	-0.003666	0.055706	-0.499512	
Contacts	-0.008565	-0.000282	-0.009801	-0.008902	-0.309932	
Yearly_equip_failure	-0.007685	0.017058	0.006632	0.056346	0.495476	
Tenure	-0.010849	0.700531	-0.073350	-0.059601	-0.023359	
MonthlyCharge	-0.000559	0.044965	-0.025080	0.696056	-0.048205	
Bandwidth_GB_Year	-0.012615	0.702102	-0.078281	-0.008758	-0.012278	
SurveyResponse	0.458796	0.032644	0.280102	0.032057	0.010777	
SurveyFixes	0.434080	0.044668	0.281116	0.018523	0.002328	
SurveyReplacements	0.400724	0.036816	0.280686	-0.010283	-0.013250	
SurveyReliability	0.145661	-0.053416	-0.566751	-0.031787	-0.003018	
SurveyOptions	-0.175523	0.067532	0.586020	0.026249	-0.015021	

SurveyRespect	0.405060	-0.013150	-0.183604	0.005018	-0.006709
SurveyCourteous	0.358304	-0.006069	-0.180813	-0.031803	0.011792
SurveyListening	0.308741	-0.016078	-0.132101	0.028892	-0.005769

	PC6	PC7	PC8	PC9	PC10	\
Population	-0.158561	0.107099	0.729421	0.046957	0.106960	
Children	-0.414608	0.311922	-0.175497	0.512950	0.339433	
Age	0.459666	0.613717	-0.025014	0.169097	-0.559523	
Income	0.386314	-0.530936	0.226829	0.699160	-0.081200	
Outage_sec_perweek	0.000618	-0.035980	0.069240	0.014744	0.030149	
Email	-0.333894	0.273176	-0.076051	0.445199	-0.126064	
Contacts	0.561927	0.231294	-0.190770	0.055106	0.699880	
Yearly_equip_failure	0.089111	0.317711	0.573607	-0.094018	0.187278	
Tenure	0.001832	0.010757	-0.002293	-0.010241	-0.020274	
MonthlyCharge	0.060946	0.036933	-0.065864	-0.025081	-0.069746	
Bandwidth_GB_Year	-0.021574	0.008498	-0.012761	-0.004181	0.005750	
SurveyResponse	0.000905	0.007292	0.003301	-0.013699	0.019661	
SurveyFixes	0.020153	0.009658	-0.015558	-0.002410	0.003798	
SurveyReplacements	-0.027931	-0.004126	-0.009081	-0.027810	-0.012818	
SurveyReliability	0.000443	-0.008382	0.006269	-0.031082	-0.004963	
SurveyOptions	0.029561	0.006610	-0.009244	-0.011694	-0.015676	
SurveyRespect	0.003374	0.008938	0.014247	0.014941	0.014856	
SurveyCourteous	0.008325	-0.019559	-0.028750	0.052351	0.044482	
SurveyListening	0.030224	-0.006635	0.055473	-0.034340	-0.060296	

	PC11	PC12	PC13	PC14	PC15	\
Population	-0.446193	-0.014801	-0.061320	0.020601	-0.016653	
Children	-0.372158	0.009287	0.036773	-0.046451	-0.014709	
Age	-0.214528	0.104919	-0.072154	0.014209	-0.003069	
Income	0.042847	-0.056917	0.002122	-0.051703	-0.002525	
Outage_sec_perweek	0.023587	0.692611	-0.117754	0.011810	-0.012325	
Email	0.578073	-0.038336	0.065146	-0.017915	-0.016313	
Contacts	0.029897	0.017236	0.038945	-0.036829	-0.003329	
Yearly_equip_failure	0.501887	-0.124765	0.030201	0.005490	-0.014937	
Tenure	0.017391	0.039345	0.000411	-0.007265	0.007949	
MonthlyCharge	-0.143934	-0.685115	0.047218	0.010463	-0.014067	
Bandwidth_GB_Year	-0.001585	-0.009198	0.011018	0.003327	0.002187	
SurveyResponse	0.008390	-0.008013	-0.071120	-0.118479	-0.047430	
SurveyFixes	-0.006463	-0.000031	-0.112110	-0.169856	-0.066158	
SurveyReplacements	0.022081	-0.011836	-0.176501	-0.247545	-0.145847	
SurveyReliability	0.009143	-0.024202	-0.180151	-0.477840	-0.444332	
SurveyOptions	-0.004841	0.041377	0.129460	0.059700	-0.209737	
SurveyRespect	0.009247	-0.003978	-0.061852	0.064204	0.757673	
SurveyCourteous	0.018457	-0.022811	-0.161773	0.807284	-0.377386	
SurveyListening	-0.069925	0.119793	0.915859	-0.025113	-0.112988	

PC16	PC17	PC18	PC19
------	------	------	------

Population	0.000779	-0.004951	-0.002278	-0.000272
Children	0.013368	0.022523	-0.005728	-0.024382
Age	-0.017409	0.010475	0.017236	0.021007
Income	-0.007898	0.012777	0.005285	0.004322
Outage_sec_perweek	-0.017055	0.010844	-0.004232	-0.000537
Email	0.006996	-0.016869	0.001158	0.002413
Contacts	-0.025880	0.020654	-0.000704	-0.000614
Yearly equip_failure	-0.000625	0.006623	-0.021210	-0.000405
Tenure	-0.014183	0.007800	0.003097	-0.704827
MonthlyCharge	0.000258	0.021226	-0.013417	-0.048510
Bandwidth_GB_Year	0.002832	-0.010185	0.010384	0.706642
SurveyResponse	0.023999	-0.239656	0.792887	-0.006345
SurveyFixes	0.068930	-0.591542	-0.573519	-0.002684
SurveyReplacements	-0.392388	0.676032	-0.177546	0.018473
SurveyReliability	0.431204	0.086404	0.017657	-0.002076
SurveyOptions	0.694894	0.261050	-0.042670	-0.004053
SurveyRespect	0.404606	0.226563	-0.063876	0.000992
SurveyCourteous	0.067671	0.065179	-0.041271	-0.006966
SurveyListening	-0.042895	0.046252	-0.042857	-0.001540

```
[50]: # Percentage of variance explained by number of components
print('Percentage of variance explained by number of components:')
for index, var in enumerate(pca.explained_variance_ratio_.cumsum()):
    print(f'- PC{index + 1}: {round(var * 100, 2)}%')
```

Percentage of variance explained by number of components:

- PC1: 15.52%
- PC2: 25.53%
- PC3: 34.14%
- PC4: 40.1%
- PC5: 45.56%
- PC6: 50.94%
- PC7: 56.29%
- PC8: 61.53%
- PC9: 66.71%
- PC10: 71.87%
- PC11: 76.99%
- PC12: 81.55%
- PC13: 85.64%
- PC14: 89.27%
- PC15: 92.39%
- PC16: 95.22%
- PC17: 97.76%
- PC18: 99.47%
- PC19: 100.0%

E1. Principal components

The scree and eigenvalue plots showed that the dataset contained 11 principal components: Pop-

ulation, Children, Age, Income, Outage_sec_perweek, Email, Contacts, Yearly_equip_failure, Tenure, MonthlyCharge, and Bandwidth_GB_Year.

E2. Identifying the principal components

To identify the principal components, I performed a Principal Component Analysis on the numeric columns that provided statistical insight from the dataset. I graphed the scree plot to look at how much variance can be explained based on the number of components. The first 11 components together explained 76.99% of the total variability. For additional information, I extracted and plotted the eigenvalues of all the components. The eigenvalue plot showed that first 11 components had eigenvalues greater than one. The eigenvalue criterion states that only components with eigenvalue greater than one should be retained because each component should explain at least one predictor's worth of variability (Larose, 2019, p. 188). The last eight components did not provide additional value when analyzing churn rate and can be reduced from the dataset.

E3. Benefits of PCA

The results of the PCA showed that the dataset had 11 principal components. Focusing on these components will help the most when attempting to explain the variability in the data. The last eight components, the survey responses, did not provide additional insight into the dataset. As such, they can be reduced from the data. This result is not a surprise since customer retention can be made up of multiple different components. There is not a single component that will predict the churn rate with certainty. However, knowing which components to focus on would save the company time and money from analyzing the unimportant variables.

Part IV: Supporting Documents

F. Panopto recording

<https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=37f4a481-e095-431a-8105-ade1013d4e8e>

G. Third-party code sources

Middleton, K. (2021). Lesson 7: Principal Component Analysis (PCA). Western Governors University. <https://my.wgu.edu/courses/course/23720006/course-material>

Python Pool. (2021, July 10). Matplotlib Boxplot With Customization in Python. [Weblog]. <https://www.pythonpool.com/matplotlib-boxplot/>

Solomon, B. (2021, June 19). Python Histogram Plotting: NumPy, Matplotlib, Pandas & Seaborn [Weblog]. <https://realpython.com/python-histograms/>

Zach, B. (2021, September 18). How to Create a Scree Plot in Python. [Weblog]. <https://www.statology.org/scree-plot-python/>

H. References

Larose, C. D. & Larose, D. T. (2019). Data Science Using Python and R. (1st ed.). Wiley.

Donthi, S. (2021). Dealing with Missing Data in Python. DataCamp. <https://app.datacamp.com/learn/courses/dealing-with-missing-data-in-python>

Nehme, A. (2021). Cleaning Data in Python. DataCamp. <https://app.datacamp.com/learn/courses/cleaning-data-in-python>