

D208 Performance Assessment Task 1

April 26, 2022

Kiet Nguyen ID: 001601720 Email: kngu179@wgu.edu

0.0.1 A. Purpose of Analysis

A1. Research Question. Which independent variables had the most significant impact on customers' MonthlyCharge?

A2. Objectives The objective of the analysis was to identify which independent variables significantly affected MonthlyCharge. Customers who were price sensitive would be more likely to churn if their monthly payments were too high. By determining which variables influenced MonthlyCharge, stakeholders could run various marketing campaigns to reduce cost for customers. This analysis could also lead to a better understanding of how the dependent variable changes based on the independent variables.

0.0.2 B. Multiple Regression Description

B1. Summarize Assumptions To perform multiple linear regression, five assumptions must be met (Bobbitt, 2021b):

1. Linear relationship between independent and dependent variable.
2. Independent variables are not highly correlated with each other.
3. Observations are independent.
4. Residuals have constant variance.
5. Residuals are normally distributed.

B2. Tool Benefits I chose Python for this project because it is a flexible and powerful programming language. The syntax was also easy to read. On top of that, Python had a great ecosystem of libraries that made data analysis tasks much easier (Rane, 2021). The libraries used in this project were:

- NumPy: high-performance numerical computation.
- Pandas: fast and flexible dataframes.
- Matplotlib and Seaborn: data visualizations.
- Statsmodels: classes and functions for different statistical models.

The environment for this project was Conda, an open-source package manager and environment management system.

B3. Multiple Regression Justification Multiple linear regression (MLR) could explain the relationship between multiple independent variables against one dependent variable (Bevans, 2020). This statistical technique could provide a more accurate representation of exactly how the dependent variable changes through a linear model. I wanted to use MLR to find out which independent variables are most significant to the dependent variable. By creating a linear regression model based on these variables, I could adjust the independent variables to see how they affect `MonthlyCharge`.

0.0.3 C. Data Preparation Process

C1. Preparation Goals The goals of data preparation included:

- Learn about the dataset and its variables.
- Explore measures of central tendency (mean, median, and mode).
- Check for missing data and handle them as necessary.
- Visualize data through univariate and bivariate plots.
- Remove highly correlated columns for MLR.

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor

churn = pd.read_csv('churn_clean.csv')
churn_cp = churn.copy(deep=True)
```

```
[2]: churn_cp.head()
```

```
[2]: CaseOrder Customer_id Interaction \
0      1      K409198 aa90260b-4141-4a24-8e36-b04ce1f4f77b
1      2      S120509 fb76459f-c047-4a9d-8af9-e0f7d4ac2524
2      3      K191035 344d114c-3736-4be5-98f7-c72c281e2d35
3      4      D90850 abfa2b40-2d43-4994-b15a-989b8c79e311
4      5      K662701 68a861fd-0d20-4e51-a587-8a90407ee574

      UID      City State      County \
0 e885b299883d4f9fb18e39c75155d990 Point Baker AK Prince of Wales-Hyder
1 f2de8bef964785f41a2959829830fb8a West Branch MI Ogemaw
2 f1784cfa9f6d92ae816197eb175d3c71 Yamhill OR Yamhill
3 dc8a365077241bb5cd5ccd305136b05e Del Mar CA San Diego
4 aabb64a116e83fdc4befc1fbab1663f9 Needville TX Fort Bend

      Zip      Lat      Lng ... MonthlyCharge Bandwidth_GB_Year Item1 \
0 99927 56.25100 -133.37571 ... 172.455519 904.536110 5
1 48661 44.32893 -84.24080 ... 242.632554 800.982766 3
2 97148 45.35589 -123.24657 ... 159.947583 2054.706961 4
3 92014 32.96687 -117.24798 ... 119.956840 2164.579412 4
```

```
4  77461  29.38012 -95.80673 ...    149.948316    271.493436    4
```

	Item2	Item3	Item4	Item5	Item6	Item7	Item8
0	5	5	3	4	4	3	4
1	4	3	3	4	3	4	4
2	4	2	4	4	3	3	3
3	4	4	2	5	4	3	3
4	4	4	3	4	4	4	5

```
[5 rows x 50 columns]
```

C2. Summary Statistics The original dataset contained 10,000 rows and 50 columns. 23 columns were dropped from the dataset because they were not relevant to the MLR. The first twelve dropped columns were customer demographic data: `CaseOrder`, `Customer_id`, `Interaction`, `UID`, `City`, `State`, `County`, `Zip`, `Lat`, `Lng`, `Area`, `TimeZone`, `Job`, `Marital`, and `PaymentMethod`. Each of these demographic columns contained multiple unique categorical values that would generate too many dummy variables. The last eight dropped columns were survey responses, `Item1` through `Item8`, that were not significant to the target dependent variable.

The prepared dataset contained 26 independent variables and one dependent variable. The target dependent variable is `MonthlyCharge`, the amount charged to the customer monthly. The mean of this variable is 172.62 with a standard deviation of 42.94.

The eleven continuous independent variables are: `Population`, `Children`, `Age`, `Income`, `Outage_sec_perweek`, `Email`, `Contacts`, `Yearly_equip_failure`, `Tenure`, `MonthlyCharge`, and `Bandwidth_GB_Year`.

The sixteen categorical independent variables are: `Gender`, `Churn`, `Techie`, `Contract`, `Port_modem`, `Tablet`, `InternetService`, `Phone`, `Multiple`, `OnlineSecurity`, `OnlineBackup`, `DeviceProtection`, `TechSupport`, `StreamingTV`, `StreamingMovies`, and `PaperlessBilling`.

```
[4]: churn_cp['MonthlyCharge'].describe()
```

```
[4]: count    10000.000000
      mean      172.624816
      std       42.943094
      min       79.978860
      25%      139.979239
      50%      167.484700
      75%      200.734725
      max       290.160419
      Name: MonthlyCharge, dtype: float64
```

```
[3]: # Drop customer demographic and survey columns that are not important to the
      ↪ regression analysis
```

```
churn_cp.drop(['CaseOrder', 'Customer_id', 'Interaction', 'UID', 'City',
↳ 'State', 'County', 'Zip', 'Lat', 'Lng', 'Area', 'TimeZone', 'Job',
↳ 'Marital', 'PaymentMethod', 'Item1', 'Item2', 'Item3', 'Item4', 'Item5',
↳ 'Item6', 'Item7', 'Item8'], axis=1, inplace=True)
```

```
churn_cp.head()
```

```
[3]:   Population  Children  Age    Income  Gender  Churn  Outage_sec_perweek  \
0          38          0   68  28561.99   Male    No          7.978323
1        10446          1   27  21704.77  Female   Yes         11.699080
2         3735          4   50   9609.57  Female    No         10.752800
3        13863          1   48  18925.23   Male    No         14.913540
4        11352          0   83  40074.19   Male   Yes          8.147417

      Email  Contacts  Yearly equip_failure  ...  OnlineSecurity  OnlineBackup  \
0        10         0                   1  ...             Yes             Yes
1        12         0                   1  ...             Yes             No
2         9         0                   1  ...             No             No
3        15         2                   0  ...             Yes             No
4        16         2                   1  ...             No             No

      DeviceProtection  TechSupport  StreamingTV  StreamingMovies  PaperlessBilling  \
0                   No           No           No             Yes             Yes
1                   No           No           Yes             Yes             Yes
2                   No           No           No             Yes             Yes
3                   No           No           Yes             No             Yes
4                   No           Yes           Yes             No             No

      Tenure  MonthlyCharge  Bandwidth_GB_Year
0    6.795513    172.455519    904.536110
1    1.156681    242.632554    800.982766
2   15.754144    159.947583   2054.706961
3   17.087227    119.956840   2164.579412
4    1.670972    149.948316    271.493436
```

```
[5 rows x 27 columns]
```

```
[5]: # Check statistics of continuous variables
churn_cp.describe()
```

```
[5]:   Population  Children  Age    Income  \
count  10000.000000  10000.0000  10000.000000  10000.000000
mean    9756.562400    2.0877   53.078400   39806.926771
std    14432.698671    2.1472   20.698882   28199.916702
min      0.000000    0.0000   18.000000    348.670000
25%     738.000000    0.0000   35.000000   19224.717500
50%    2910.500000    1.0000   53.000000   33170.605000
```

75%	13168.000000	3.0000	71.000000	53246.170000
max	111850.000000	10.0000	89.000000	258900.700000

	Outage_sec_perweek	Email	Contacts	Yearly_equip_failure \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	10.001848	12.016000	0.994200	0.398000
std	2.976019	3.025898	0.988466	0.635953
min	0.099747	1.000000	0.000000	0.000000
25%	8.018214	10.000000	0.000000	0.000000
50%	10.018560	12.000000	1.000000	0.000000
75%	11.969485	14.000000	2.000000	1.000000
max	21.207230	23.000000	7.000000	6.000000

	Tenure	MonthlyCharge	Bandwidth_GB_Year
count	10000.000000	10000.000000	10000.000000
mean	34.526188	172.624816	3392.341550
std	26.443063	42.943094	2185.294852
min	1.000259	79.978860	155.506715
25%	7.917694	139.979239	1236.470827
50%	35.430507	167.484700	3279.536903
75%	61.479795	200.734725	5586.141370
max	71.999280	290.160419	7158.981530

```
[6]: # Check statistics of categorical variable
churn_cp.describe(include=object)
```

```
[6]:
```

	Gender	Churn	Techie	Contract	Port_modem	Tablet \
count	10000	10000	10000	10000	10000	10000
unique	3	2	2	3	2	2
top	Female	No	No	Month-to-month	No	No
freq	5025	7350	8321	5456	5166	7009

	InternetService	Phone	Multiple	OnlineSecurity	OnlineBackup \
count	10000	10000	10000	10000	10000
unique	3	2	2	2	2
top	Fiber Optic	Yes	No	No	No
freq	4408	9067	5392	6424	5494

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies \
count	10000	10000	10000	10000
unique	2	2	2	2
top	No	No	No	No
freq	5614	6250	5071	5110

	PaperlessBilling
count	10000
unique	2

top	Yes
freq	5882

C3. Preparation Steps The steps used to prepare the data are:

1. Get an overview of the imported data.
2. Drop columns unnecessary for the MLR.
3. Check statistics of continuous and categorical variables.
4. Check for any duplicate or missing values.
5. Plot univariate and bivariate plots.
6. Encode categorical variables into numerical values.
7. Run a heatmap and drop highly correlated columns.
8. Calculate VIF and drop columns with $VIF > 5$.

```
[7]: # Check if there is any duplicates
churn_cp.duplicated().any()
```

[7]: False

```
[8]: # Check if there is any missing values
churn_cp.isnull().values.any()
```

[8]: False

C4. Generate Visualizations

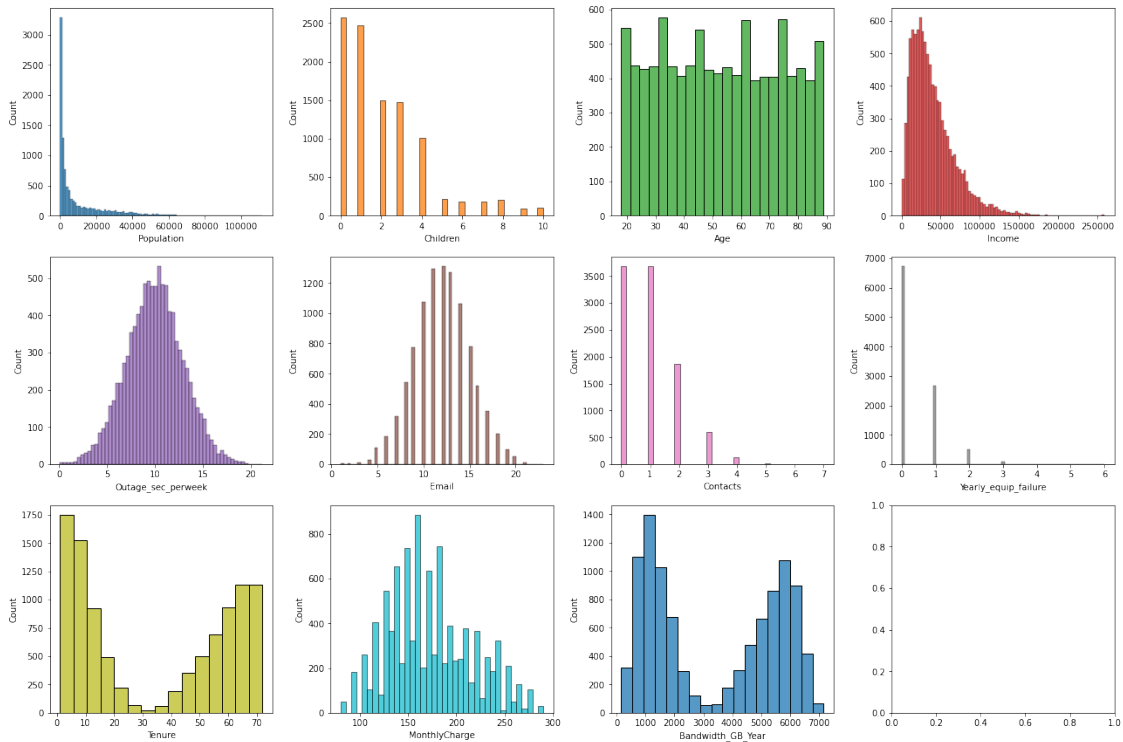
```
[9]: # Plot histograms for continuous columns
continuous = churn_cp.select_dtypes(include='number').columns.tolist()

fig, axes = plt.subplots(3, 4, figsize=(18, 12))
palette1 = sns.color_palette()

# Adapted from seaborn documentation (Waskom, 2022)
# https://seaborn.pydata.org/tutorial/color_palettes.html
x = 0
y = 0
color = 0
for col in continuous:
    if color == 10:
        color = 0
    if y == 4:
        x += 1
        y = 0
    sns.histplot(ax=axes[x, y], data=churn_cp[col], color=palette1[color])
    color += 1
    y += 1

plt.tight_layout()
```

```
plt.show()
```

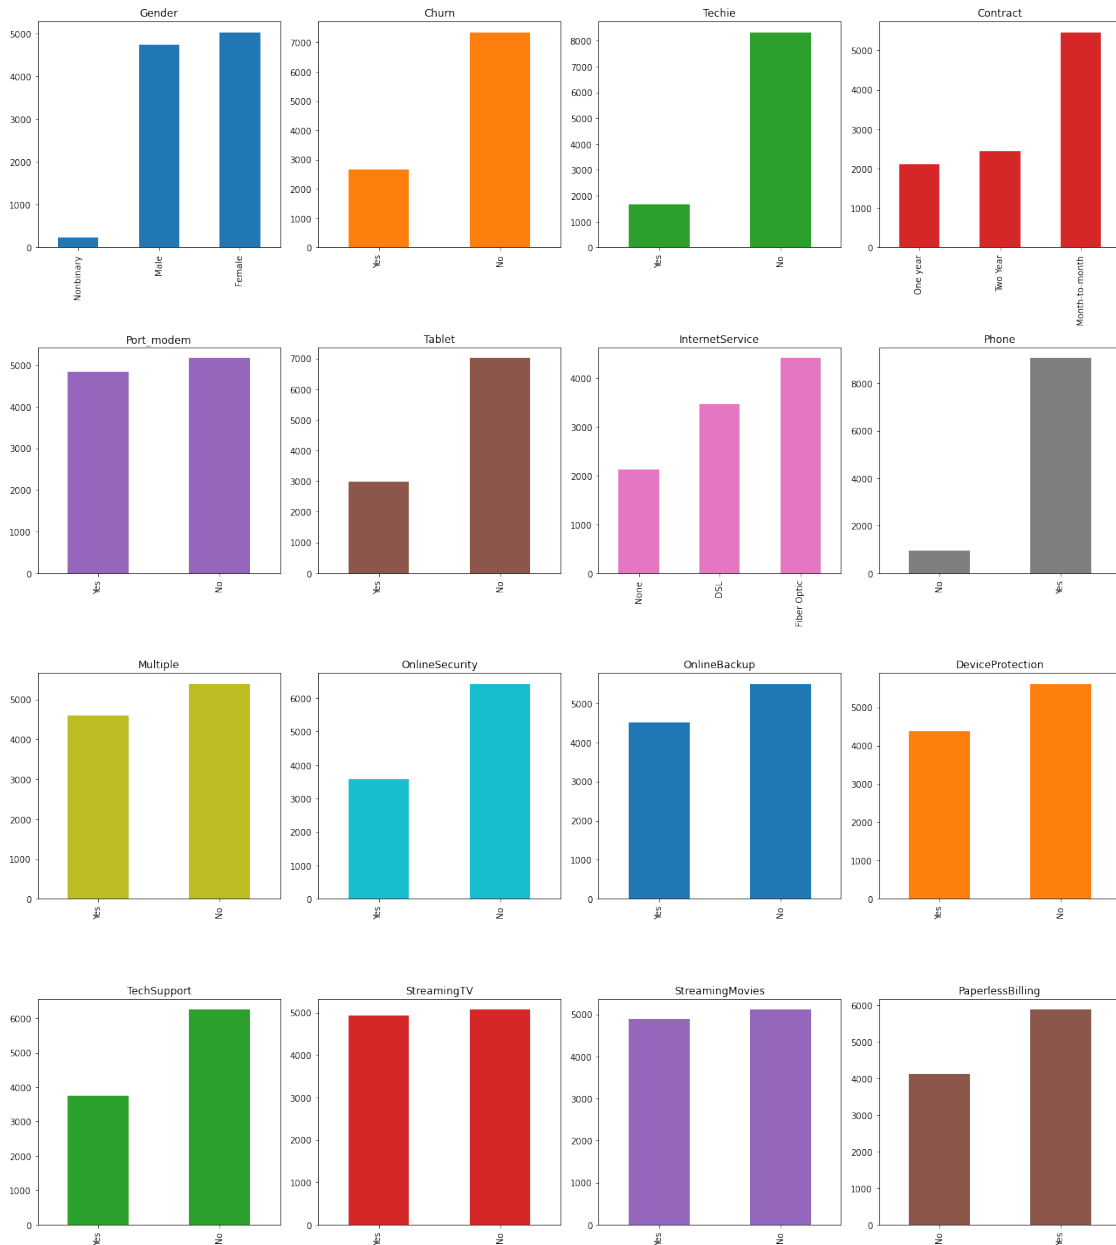


```
[10]: # Plot barplots for categorical columns
categorical = churn_cp.select_dtypes(include='object').columns.tolist()

fig, axes = plt.subplots(4, 4, figsize=(18, 20))
palette2 = sns.color_palette('tab10')

x = 0
y = 0
color = 0
for col in categorical:
    if color == 10:
        color = 0
    if y == 4:
        x += 1
        y = 0
    count = churn_cp[col].value_counts().sort_values()
    axes[x, y].set_title(col)
    count.plot.bar(ax=axes[x, y], color=palette2[color])
    color += 1
    y += 1
```

```
plt.tight_layout()
plt.show()
```



```
[11]: # Generate bivariate pairplots
g = sns.PairGrid(churn_cp)
g.map_upper(sns.scatterplot, color='#333C83')
g.map_diag(sns.histplot, color='#F24A72')
g.map_lower(sns.scatterplot, color='#006778')
plt.show()
```




```
[12]: # Convert categorical columns into dummy variables
churn_dmy = pd.get_dummies(churn_cp, drop_first=True)

# Replace space in column names with underscore
churn_dmy.columns = churn_dmy.columns.str.replace(' ', '_')

churn_dmy.describe()
```

```
[12]:
```

	Population	Children	Age	Income \
count	10000.000000	10000.0000	10000.000000	10000.000000
mean	9756.562400	2.0877	53.078400	39806.926771
std	14432.698671	2.1472	20.698882	28199.916702
min	0.000000	0.0000	18.000000	348.670000

25%	738.000000	0.0000	35.000000	19224.717500
50%	2910.500000	1.0000	53.000000	33170.605000
75%	13168.000000	3.0000	71.000000	53246.170000
max	111850.000000	10.0000	89.000000	258900.700000

	Outage_sec_perweek	Email	Contacts	Yearly_equip_failure \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	10.001848	12.016000	0.994200	0.398000
std	2.976019	3.025898	0.988466	0.635953
min	0.099747	1.000000	0.000000	0.000000
25%	8.018214	10.000000	0.000000	0.000000
50%	10.018560	12.000000	1.000000	0.000000
75%	11.969485	14.000000	2.000000	1.000000
max	21.207230	23.000000	7.000000	6.000000

	Tenure	MonthlyCharge ...	InternetService_None	Phone_Yes \
count	10000.000000	10000.000000 ...	10000.000000	10000.000000
mean	34.526188	172.624816 ...	0.212900	0.906700
std	26.443063	42.943094 ...	0.409378	0.290867
min	1.000259	79.978860 ...	0.000000	0.000000
25%	7.917694	139.979239 ...	0.000000	1.000000
50%	35.430507	167.484700 ...	0.000000	1.000000
75%	61.479795	200.734725 ...	0.000000	1.000000
max	71.999280	290.160419 ...	1.000000	1.000000

	Multiple_Yes	OnlineSecurity_Yes	OnlineBackup_Yes \
count	10000.000000	10000.000000	10000.000000
mean	0.460800	0.357600	0.450600
std	0.498486	0.479317	0.497579
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000

	DeviceProtection_Yes	TechSupport_Yes	StreamingTV_Yes \
count	10000.000000	10000.000000	10000.000000
mean	0.438600	0.375000	0.492900
std	0.496241	0.484147	0.499975
min	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000
75%	1.000000	1.000000	1.000000
max	1.000000	1.000000	1.000000

	StreamingMovies_Yes	PaperlessBilling_Yes
count	10000.000000	10000.000000

```

mean      0.489000      0.588200
std       0.499904      0.492184
min       0.000000      0.000000
25%       0.000000      0.000000
50%       0.000000      1.000000
75%       1.000000      1.000000
max       1.000000      1.000000

```


[8 rows x 30 columns]

```

[13]: # Run a correlation matrix
churn_corr = churn_dmy.corr()

# Plot a heatmap of the correlations
plt.figure(figsize=(26, 14))
sns.heatmap(churn_corr, annot=True, xticklabels=churn_corr.columns,
            yticklabels=churn_corr.columns, cmap='coolwarm')
plt.show()

```



```

[14]: # Drop highly correlated columns
churn_dmy.drop(['Bandwidth_GB_Year', 'Tenure'], axis=1, inplace=True)

churn_dmy.head()

```

```
[14]: Population  Children  Age    Income  Outage_sec_perweek  Email  Contacts  \
0          38          0   68  28561.99          7.978323      10          0
1       10446          1   27  21704.77          11.699080      12          0
2        3735          4   50   9609.57          10.752800       9          0
3       13863          1   48  18925.23          14.913540      15          2
4       11352          0   83  40074.19           8.147417      16          2

    Yearly_equip_failure  MonthlyCharge  Gender_Male  ...  \
0                1      172.455519          1  ...
1                1      242.632554          0  ...
2                1      159.947583          0  ...
3                0      119.956840          1  ...
4                1      149.948316          1  ...

    InternetService_None  Phone_Yes  Multiple_Yes  OnlineSecurity_Yes  \
0                0          1          0          1
1                0          1          1          1
2                0          1          1          0
3                0          1          0          1
4                0          0          0          0

    OnlineBackup_Yes  DeviceProtection_Yes  TechSupport_Yes  StreamingTV_Yes  \
0                1                0          0          0
1                0                0          0          1
2                0                0          0          0
3                0                0          0          1
4                0                0          1          1

    StreamingMovies_Yes  PaperlessBilling_Yes
0                1                1
1                1                1
2                1                1
3                0                1
4                0                0
```

[5 rows x 28 columns]

```
[15]: # Adapted from Detecting Multicollinearity with VIF (GeeksforGeeks, 2020)
# https://www.geeksforgeeks.org/detecting-multicollinearity-with-vif-python/

# Create independent variables set
X = churn_dmy.drop('MonthlyCharge', axis=1)

# VIF dataframe
vif = pd.DataFrame()
vif['Variables'] = X.columns
```

```
# Calculating VIF for each variable
vif['VIF'] = [variance_inflation_factor(X.values, i) for i in range(len(X.
    ↪columns)))]

vif
```

```
[15]:
```

	Variables	VIF
0	Population	1.452552
1	Children	1.919980
2	Age	6.799319
3	Income	2.881564
4	Outage_sec_perweek	10.106537
5	Email	12.448311
6	Contacts	1.988874
7	Yearly_equip_failure	1.384885
8	Gender_Male	1.925507
9	Gender_Nonbinary	1.046806
10	Churn_Yes	1.811674
11	Techie_Yes	1.205180
12	Contract_One_year	1.455095
13	Contract_Two_Year	1.542198
14	Port_modem_Yes	1.915201
15	Tablet_Yes	1.424275
16	InternetService_Fiber_Optic	2.237575
17	InternetService_None	1.598786
18	Phone_Yes	8.980119
19	Multiple_Yes	1.879134
20	OnlineSecurity_Yes	1.544828
21	OnlineBackup_Yes	1.804321
22	DeviceProtection_Yes	1.777888
23	TechSupport_Yes	1.590844
24	StreamingTV_Yes	2.097631
25	StreamingMovies_Yes	2.158125
26	PaperlessBilling_Yes	2.372637

```
[16]: # Get a list of variables with VIF > 5
high_vif = vif[vif['VIF'] > 5]
high_vif = high_vif['Variables'].tolist()

# Drop columns with high VIF
churn_dmy.drop(high_vif, axis=1, inplace=True)

churn_dmy.head()
```

```
[16]:
```

	Population	Children	Income	Contacts	Yearly_equip_failure	\
0	38	0	28561.99	0	1	
1	10446	1	21704.77	0	1	

2	3735	4	9609.57	0	1
3	13863	1	18925.23	2	0
4	11352	0	40074.19	2	1

	MonthlyCharge	Gender_Male	Gender_Nonbinary	Churn_Yes	Techie_Yes	...	\
0	172.455519	1	0	0	0	...	
1	242.632554	0	0	1	1	...	
2	159.947583	0	0	0	1	...	
3	119.956840	1	0	0	1	...	
4	149.948316	1	0	1	0	...	

	InternetService_Fiber_Optic	InternetService_None	Multiple_Yes	\
0	1	0	0	
1	1	0	1	
2	0	0	1	
3	0	0	0	
4	1	0	0	

	OnlineSecurity_Yes	OnlineBackup_Yes	DeviceProtection_Yes	\
0	1	1	0	
1	1	0	0	
2	0	0	0	
3	1	0	0	
4	0	0	0	

	TechSupport_Yes	StreamingTV_Yes	StreamingMovies_Yes	PaperlessBilling_Yes
0	0	0	1	1
1	0	1	1	1
2	0	0	1	1
3	0	1	0	1
4	1	1	0	0

[5 rows x 24 columns]

C5. Copy of Prepared Data

```
[17]: # Create a copy of the prepared data
churn_dmy.to_csv('churn_prepare.csv', index=False)
```

0.0.4 D. Compare Initial and Reduced Models

D1. Construct Initial Model

```
[18]: Y = churn_dmy['MonthlyCharge']
X = churn_dmy.drop('MonthlyCharge', axis=1)

X = sm.add_constant(X)
model = sm.OLS(Y, X).fit()
```

```
print(model.summary())
```

OLS Regression Results

```
=====
Dep. Variable:          MonthlyCharge    R-squared:                0.959
Model:                  OLS              Adj. R-squared:          0.959
Method:                 Least Squares    F-statistic:             1.006e+04
Date:                   Tue, 26 Apr 2022  Prob (F-statistic):       0.00
Time:                   15:27:29          Log-Likelihood:          -35859.
No. Observations:      10000             AIC:                    7.177e+04
Df Residuals:          9976              BIC:                    7.194e+04
Df Model:               23
Covariance Type:       nonrobust
=====
```

```
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                83.6796      0.372    225.111     0.000
82.951    84.408
Population           1.089e-06   6.06e-06     0.180     0.857
-1.08e-05    1.3e-05
Children              0.0062      0.041     0.151     0.880
-0.074      0.086
Income              3.311e-06   3.1e-06     1.067     0.286
-2.77e-06   9.39e-06
Contacts            -0.0462      0.089    -0.522     0.602
-0.220      0.127
Yearly_equip_failure -0.0985      0.138    -0.716     0.474
-0.368      0.171
Gender_Male         -0.2929      0.177    -1.652     0.099
-0.640      0.055
Gender_Nonbinary    -0.8302      0.589    -1.410     0.159
-1.985      0.324
Churn_Yes            2.5025      0.229    10.946     0.000
2.054      2.951
Techie_Yes           0.2383      0.235     1.015     0.310
-0.222      0.698
Contract_One_year    0.8026      0.231     3.474     0.001
0.350      1.255
Contract_Two_Year    0.7724      0.220     3.503     0.000
0.340      1.205
Port_modem_Yes      -0.2277      0.175    -1.301     0.193
-0.571      0.115
Tablet_Yes          -0.1759      0.191    -0.920     0.358
-0.551      0.199
=====
```

InternetService_Fiber_Optic	20.0524	0.200	100.419	0.000
19.661 20.444				
InternetService_None	-12.6692	0.242	-52.409	0.000
-13.143 -12.195				
Multiple_Yes	32.2933	0.178	181.711	0.000
31.945 32.642				
OnlineSecurity_Yes	2.7101	0.183	14.837	0.000
2.352 3.068				
OnlineBackup_Yes	22.4526	0.176	127.456	0.000
22.107 22.798				
DeviceProtection_Yes	12.4009	0.177	70.193	0.000
12.055 12.747				
TechSupport_Yes	12.5063	0.181	69.165	0.000
12.152 12.861				
StreamingTV_Yes	41.6656	0.181	229.564	0.000
41.310 42.021				
StreamingMovies_Yes	51.6903	0.184	280.210	0.000
51.329 52.052				
PaperlessBilling_Yes	0.1430	0.178	0.804	0.421
-0.206 0.492				
=====				
Omnibus:	35663.845	Durbin-Watson:	1.997	
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1567.455	
Skew:	0.021	Prob(JB):	0.00	
Kurtosis:	1.061	Cond. No.	3.34e+05	
=====				

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.34e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
[19]: # Residuals of the initial model
      model.resid
```

```
[19]: 0      -8.374701
      1       7.842365
      2     -8.603571
      3     -8.950703
      4     -9.891661
      ...
     9995    9.738315
     9996   -8.382081
     9997    8.484229
     9998   10.235932
     9999   -7.420766
```


Length: 10000, dtype: float64

D2. Variable Selection Procedure

```
[20]: # Get a list of variables with p > 0.05
high_p = []
for col in churn_dmy.columns:
    if col == 'MonthlyCharge':
        pass
    else:
        p_val = model.pvalues[col]
        if p_val > 0.05:
            print(f'{col}: {p_val:.2f}')
            high_p.append(col)
```

Population: 0.86
Children: 0.88
Income: 0.29
Contacts: 0.60
Yearly_equip_failure: 0.47
Gender_Male: 0.10
Gender_Nonbinary: 0.16
Techie_Yes: 0.31
Port_modem_Yes: 0.19
Tablet_Yes: 0.36
PaperlessBilling_Yes: 0.42

```
[21]: # Perform backward selection by dropping columns with high p-values
churn_dmy.drop(high_p, axis=1, inplace=True)

churn_dmy.head()
```

```
[21]:
```

	MonthlyCharge	Churn_Yes	Contract_One_year	Contract_Two_Year	\
0	172.455519	0	1	0	
1	242.632554	1	0	0	
2	159.947583	0	0	1	
3	119.956840	0	0	1	
4	149.948316	1	0	0	

	InternetService_Fiber_Optic	InternetService_None	Multiple_Yes	\
0	1	0	0	
1	1	0	1	
2	0	0	1	
3	0	0	0	
4	1	0	0	

	OnlineSecurity_Yes	OnlineBackup_Yes	DeviceProtection_Yes	\
0	1	1	0	

1	1	0	0
2	0	0	0
3	1	0	0
4	0	0	0

	TechSupport_Yes	StreamingTV_Yes	StreamingMovies_Yes
0	0	0	1
1	0	1	1
2	0	0	1
3	0	1	0
4	1	1	0

```
[22]: # Drop variables that have coefficients < 3
churn_dmy.drop(['Churn_Yes', 'Contract_One_year', 'Contract_Two_Year',
↳ 'OnlineSecurity_Yes'], axis=1, inplace=True)
```

D3. Reduced Regression Model

```
[23]: Y = churn_dmy['MonthlyCharge']
X = churn_dmy.drop('MonthlyCharge', axis=1)

X = sm.add_constant(X)
reduced_model = sm.OLS(Y, X).fit()
print(reduced_model.summary())
```

```

                                OLS Regression Results
=====
Dep. Variable:          MonthlyCharge    R-squared:                0.957
Model:                  OLS             Adj. R-squared:          0.957
Method:                 Least Squares    F-statistic:             2.793e+04
Date:                  Tue, 26 Apr 2022  Prob (F-statistic):       0.00
Time:                  15:28:50          Log-Likelihood:          -36031.
No. Observations:      10000            AIC:                   7.208e+04
Df Residuals:          9991             BIC:                   7.214e+04
Df Model:               8
Covariance Type:       nonrobust
=====
=====
                                coef    std err          t      P>|t|
-----
[0.025    0.975]
-----
const                84.8067      0.251    338.393     0.000
84.315    85.298
InternetService_Fiber_Optic  19.8588      0.202    98.365     0.000
19.463    20.254
InternetService_None   -12.9291      0.245   -52.810     0.000
-13.409   -12.449

```

Multiple_Yes	32.5912	0.178	182.735	0.000
32.242 32.941				
OnlineBackup_Yes	22.5885	0.179	126.453	0.000
22.238 22.939				
DeviceProtection_Yes	12.5566	0.179	70.089	0.000
12.205 12.908				
TechSupport_Yes	12.5919	0.184	68.557	0.000
12.232 12.952				
StreamingTV_Yes	42.1879	0.178	237.303	0.000
41.839 42.536				
StreamingMovies_Yes	52.3453	0.178	294.327	0.000
51.997 52.694				

Omnibus:	36028.412	Durbin-Watson:	1.992
Prob(Omnibus):	0.000	Jarque-Bera (JB):	1541.892
Skew:	0.023	Prob(JB):	0.00
Kurtosis:	1.077	Cond. No.	6.04

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
[24]: # Residuals of the reduced model
      reduced_model.resid
```

```
[24]: 0      -7.143755
      1      10.842737
      2     -9.795579
      3     -7.037777
      4     -9.496975
      ...
      9995     7.436309
      9996    -7.108647
      9997    10.128885
      9998     8.242267
      9999    -7.263100
      Length: 10000, dtype: float64
```

0.0.5 E. Analyze Dataset

E1. Explain Analysis Process Variable Selection Logic

The process of backward selection removed the variables with high p-values in the initial model. The stopping rule is the p-value threshold of 0.05 (JMP, 2019). Another factor to consider with these variables was that their coefficients were extremely small. Take **Population** in the initial model for example, the p-value of this variable was 0.857 and its coefficient was 1.089e-06. Removing this variable had minimal effect on the model. Other variables with high p-value also tend to have

coefficients that were lower than ± 1 . These variables do not significantly affect the accuracy of the model.

Model Evaluation Metric

```
[25]: # Check values of R-squared
initial_r2 = model.rsquared_adj
reduced_r2 = reduced_model.rsquared_adj

print(f'Adjusted R-squared:\n- Initial: {initial_r2:.3f}\n- Reduced: \n- {reduced_r2:.3f}')
```

Adjusted R-squared:

- Initial: 0.959
- Reduced: 0.957

The first metric to look at is the Adjusted R-squared. This score is a measure of how much variation in the dependent variable is explained by only the independent variables that actually affect the dependent variable (Glen, 2021). The score on the reduced model was 0.957 while the initial model was 0.959. That was a minimal loss in explained variation despite reducing the number of independent variables from 25 to 8.

```
[26]: # Check values of RMSE
initial_rmse = np.sqrt(model.mse_resid)
reduced_rmse = np.sqrt(reduced_model.mse_resid)

print(f'RMSE:\n- Initial: {initial_rmse:.2f}\n- Reduced: {reduced_rmse:.2f}')
```

RMSE:

- Initial: 8.74
- Reduced: 8.89

The second metric is the Root Mean Square Error (RMSE). This value is the measure of accuracy for the regression model. It is based on the average distance between the observed and predicted data values (Bobbit, 2021a). The initial model RMSE was 8.74 while the reduced model was 8.88. The slight increase in RMSE was due to the reduction of predictor variables. Overall it should not affect the accuracy of the reduced model.

Residual Plot

The scale-location plot below showed the square root of the standardized residuals versus the fitted values. This plot indicates whether the size of the residuals get bigger or smaller (Broeck, 2022). Since the residuals appear to be randomly scattered, there did not seem to be a problem with heteroscedasticity.

```
[27]: # Adapted from Assessing Model Fit (Broeck, 2022)
# https://app.datacamp.com/learn/courses/
# introduction-to-regression-with-statsmodels-in-python


# Create a scale-location plot
```

```

reduced_model_norm_resid = reduced_model.get_influence().
    ↪ resid_studentized_internal
reduced_model_norm_resid_sqrt = np.sqrt(np.abs(reduced_model_norm_resid))
sns.regplot(x=reduced_model.fittedvalues, y=reduced_model_norm_resid_sqrt,
    ↪ lowess=True, marker='+', line_kws={'color': 'red'}, scatter_kws={'s': 15})

plt.xlabel('Fitted Values')
plt.ylabel('Square Root of Standardized Residuals')
plt.show()

```



E2. Analysis Output The output and calculations of both the model and the reduced model are found in sections D1 through D3. The initial model's residuals can be found at the end of section D1. The reduced model's residuals can be found at the end of section D3.

The reduced model's predictions are below:

```

[28]: predictions = reduced_model.get_prediction()
      predictions.summary_frame()

```

```

[28]:      mean    mean_se  mean_ci_lower  mean_ci_upper  obs_ci_lower  \
0      179.599274  0.246765    179.115564    180.082984    162.172045
1      231.789817  0.245958    231.307691    232.271944    214.362632
2      169.743162  0.255481    169.242368    170.243956    152.315450
3      126.994617  0.250331    126.503918    127.485316    109.567192
4      159.445291  0.256945    158.941628    159.948954    142.017497

```

```

...
9995  152.543091  0.269429  152.014957  153.071226  135.114573
9996  214.589747  0.258382  214.083266  215.096228  197.161871
9997  159.845215  0.250246  159.354681  160.335748  142.417795
9998  244.381733  0.265728  243.860853  244.902614  226.953433
9999  224.747100  0.256063  224.245165  225.249036  207.319356

```

```

      obs_ci_upper
0      197.026503
1      249.217003
2      187.170873
3      144.422042
4      176.873086

```

```

...
9995  169.971610
9996  232.017623
9997  177.272634
9998  261.810034
9999  242.174845

```

[10000 rows x 6 columns]

E3. Code for MLR The code that ran the MLR could be found in section D1 for the initial model and D3 for the reduced model.

0.0.6 F. Summarize Findings

F1. Results of Analysis Regression Equation

```

[29]: # Lists of variables and their coefficients
cols = churn_dmy.columns.tolist()
coefs = reduced_model.params.tolist()

# Create equation string for the reduced model
equation = f'y = {coefs[0]:.2f}'
for col, coef in zip(cols[1:], coefs[1:]):
    equation += f' + ({coef:.2f} * {col})'

print(f'Regression Equation:\n{equation}')

```

Regression Equation:

```

y = 84.81 + (19.86 * InternetService_Fiber_Optic) + (-12.93 *
InternetService_None) + (32.59 * Multiple_Yes) + (22.59 * OnlineBackup_Yes) +
(12.56 * DeviceProtection_Yes) + (12.59 * TechSupport_Yes) + (42.19 *
StreamingTV_Yes) + (52.35 * StreamingMovies_Yes)

```

Interpretation of Coefficients

The majority of the coefficients in the model are positive. This is in line with the fact that having additional services and add-ons cost more money. The only negative coefficient was not having an internet service. This makes sense because the highest coefficients were streaming movies, streaming TV, and having multiple lines. All of those high cost services required the internet.

Model Significance

All of the independent variables in the equation are statistically significant. This means that the relationships between these variables and **MonthlyCharge** are statistically significant. Changes in these independent variables are associated with changes in the dependent variable when applied to the population (Frost, 2021). At the same time, there could be other factors that could influence either the coefficients or the p-values of these variables. The original dataset contained many independent variables that were reduced for the analysis. This would be an important case where having domain expertise will help the analysis identify the important variables.

Limitations of Analysis

There were multiple limitations to this analysis. The first was mentioned above, where reducing too many variables could have affected the accuracy of the final result. The second was that creating dummy variables for the categorical columns could have created some redundancy. The data generated from these dummy variables could have created some bias in the model, leading to some inaccuracy. These dummy variables also made it difficult to plot residuals versus fitted values because their values were either zero or one, not continuous.

F2. Recommendations Based on the model, the variables with the highest coefficients were streaming movies, streaming TV, and having multiple lines. Given that the mean **MonthlyCharge** was 172.62, customers who had all of mentioned additional services and add-ons would be paying significantly more than the average. This could lead to more churn for customers sensitive to high prices. To counteract this, stakeholders had multiple options:

1. Lowering the basic cost of service. This basic cost was the constant of 84.81, the price that customers paid without any of the additional variables involved. This constant alone is almost half of the mean monthly payment.
2. Lowering the costs of services and add-ons. The model indicated that having multiple lines and streaming services were expensive. Lowering the costs of these services could lead to better customer retention. The loss in profit could be made up by spreading it over the additional tenure of the customers.
3. Creating packages for customers who had these services together. Customers who had multiple services together would be less likely to leave if they are paying less by bundling these services together.

0.0.7 G. Panopto Recording

Link: <https://wgu.hosted.panopto.com/Panopto/Pages/Viewer.aspx?id=2cbfe92e-1314-4338-a99b-ae830088f19a>

0.0.8 H. Third-Party Code

Detecting multicollinearity with VIF - python. GeeksforGeeks. (2020, August 29). Retrieved April 23, 2022, from <https://www.geeksforgeeks.org/detecting-multicollinearity-with-vif-python/>

Example of multiple linear regression in python. Data to Fish. (2021, May 17). Retrieved April 22, 2022, from <https://datatofish.com/multiple-linear-regression-python/>

Middleton, K. (2021). Lesson 7: Principal Component Analysis (PCA). Western Governors University. <https://my.wgu.edu/courses/course/23720006/course-material>

Waskom, M. (2022). Choosing color palettes. Choosing color palettes - seaborn 0.11.2 documentation. Retrieved April 25, 2022, from https://seaborn.pydata.org/tutorial/color_palettes.html

0.0.9 I. References

Bevans, R. (2020, October 26). An introduction to multiple linear regression. Scribbr. Retrieved April 20, 2022, from <https://www.scribbr.com/statistics/multiple-linear-regression/>

Bobbitt, Z. (2021a, May 10). How to interpret root mean square error (RMSE). Statology. Retrieved April 24, 2022, from <https://www.statology.org/how-to-interpret-rmse/>

Broeck, M. V. den. (2022). Introduction to Regression with statsmodels in Python. DataCamp. Retrieved April 25, 2022, from <https://app.datacamp.com/learn/courses/introduction-to-regression-with-statsmodels-in-python>

Bobbitt, Z. (2021b, November 16). The five assumptions of multiple linear regression. Statology. Retrieved April 20, 2022, from <https://www.statology.org/multiple-linear-regression-assumptions/>

Frost, J. (2021). How to Interpret P-values and Coefficients in Regression Analysis. Statistics By Jim. Retrieved April 26, 2022, from <https://statisticsbyjim.com/>

Glen, S. (2021, June 7). Adjusted R2 / adjusted R-squared: What is it used for? Statistics How To. Retrieved April 24, 2022, from <https://www.statisticshowto.com/probability-and-statistics/statistics-definitions/adjusted-r2/>

Rane, Z. (2021, August 19). 10 compelling reasons to learn Python for data science. Medium. Retrieved April 20, 2022, from <https://towardsdatascience.com/10-compelling-reasons-to-learn-python-for-data-science-fa31160321cb#1faf>

Variable selection in multiple regression. JMP. (2019, January 28). Retrieved April 24, 2022, from https://www.jmp.com/en_in/statistics-knowledge-portal/what-is-multiple-regression/variable-selection.html